

From Vibe Coding to Agentic Harness Engineering (agy & claude)

Two-Phase Intensive Engineering Course | Author: Alejandro Kaspar - AI Forward Deployed Engineer (CC BY 4.0 / Apache 2.0)

Phase 1 (Modules 1–4): Vibe Coding Foundations & Workflows

Mastering Daily Execution in Google Antigravity (agy) & Claude Code (claude):

- **Module 1 (Lab 01):** Language-to-code foundations, the 2022–2026 evolution, Andrew Ng’s 4-Pillar AI Engineering Skills Map, and the Ladder of Earned Autonomy (`prompt_steerer.py`).
- **Module 2 (Lab 02):** Inside the 3-Layer Agent Stack, `agy` 101, `claude` 101, Swiss Army Knife tool primitives, and Wireshark-style harness trace inspection (`wire_trace_inspector.py`).
- **Module 3 (Lab 03):** The 5-step `RePPIT` loop (`Research` → `Propose` → `Plan` → `Implement` → `Verify`), orthogonal proposals + context reset, and reflection (`reppit_orchestrator.py`).
- **Module 4 (Lab 04):** Stateless context window physics, the 40% “Dumb Zone”, output-redirection backpressure, and Frequent Intentional Compaction (`context_compactor.py`).

Phase 2 (Modules 5–8): Harness Engineering & Agentic Loops

Architecting Agent = Model + Harness for Production Codebases:

- **Module 5 (Lab 05):** Cybernetic 2x2 Control Matrix, the 150-Instruction Cliff, the 60-line `CLAUDE.md`/`GEMINI.md` map rule, and cross-session auto-memory (`context_compiler.py`).
- **Module 6 (Lab 06):** `agentskills.io` 3-level progressive disclosure, `context-hub` (`chub`) self-annotating docs, `SkillsBench` (+16.6pp), and Meta-MCP “Code Mode” (>90% savings).
- **Module 7 (Lab 07):** Subagents as context firewalls, 7-phase `feature-dev`, CS146S PR triage (`Must Fix` / `Recommended` / `Consider`), and Karpathy’s `llm-council` (`subagent_council.py`).
- **Module 8 (Lab 08):** 4-Tier “Governed by Design” hooks, remediation linters, Karpathy’s `autoresearch` ratchet, why dark factories fail, and Recursive Self-Improvement (`autoresearch_harness.py`).

1-to-1 Dual-Branch Parity (agy & claude) + 8 Self-Diagnosing Python Labs (`./self_diagnose_all.sh` in < 1s)

- **Google Antigravity (agy) Branch:** `GEMINI.md` / `AGENTS.md`, `.agents/rules/`, `.agents/skills/`, `.agents/agents/`, `.agents/workflows/`, `.agents/plugins/`, and `KNOWLEDGE.md`.
- **Anthropic Claude Code (claude) Branch:** `CLAUDE.md`, `.claude/rules/`, `.claude/skills/`, `.claude/agents/`, `.claude-plugin/plugin.json`, and `MEMORY.md`.
- **Open-Source Licensing:** Dual-licensed under **CC BY 4.0** and **Apache 2.0** with zero external package dependencies required to run all 8 labs.

Why Language Produces Code & Actions: Fowler, Wittgenstein & Ng

Unmesh Joshi & Martin Fowler (*What Is Code?*), Wittgenstein's Two Theories of Language, and Andrew Ng's SE Steering Rule

1. Unmesh Joshi & Martin Fowler

What Is Code? (martinfowler.com):

- Code is not merely syntax for a machine; it is an explicit **externalized conceptual model**.
- Passive “vibe reviewing” creates massive **Cognitive Debt** (synonym drift, untyped dicts, silent assumptions).
- As syntax becomes cheap, **discovering precise domain vocabulary** and invariants becomes the core engineering bottleneck.

2. Wittgenstein's Two Theories

Why Natural Language Coordinates Code:

- **Early *Tractatus* (Picture Theory)**: Formal 1-to-1 logical structure—analogue to compilers, static types, and deterministic unit tests.
- **Late *Investigations* (Language-as-Action)**: Words are active tools in a shared workspace ([arXiv:1606.02447](https://arxiv.org/abs/1606.02447), [arXiv:2302.01570](https://arxiv.org/abs/2302.01570)).
- Collaborative prompting aligns intent so the agent emits formal, test-verified code.

3. Andrew Ng's SE Steering Rule

Why Fundamentals Beat Raw Vibes:


- A novice who vibe codes without **Software Engineering Fundamentals** lets the agent make poor default tradeoffs.
- Without explicit constraints, agents ignore **latency, consistency, idempotency, reliability, maintainability, and blast radius**.
- Precise architectural vocabulary is the steering wheel of an AI coding agent.

The Conceptual Synthesis: Conversational Alignment → Architectural Constraints → Deterministic Verification

- **Naive Vibe Prompt**: "Build a webhook payment handler" → Produces non-idempotent float math, unbounded retries, and SQL injection vectors.
- **SE-Steered Prompt**: "Implement an idempotent webhook handler using integer minor units, p99 < 50ms bounded retries, and strict schema validation."

The 2022–2026 Architectural Timeline, 17-Day Flagship Pace & RSI

Chronological evolution + SearchIntel’s quarterly release surge, inter-flagship gap compression (37d → 17d), and METR curve



Andrew Ng’s 4-Pillar AI Engineering Skills Map & Earned Autonomy

Mapping DeepLearning.AI’s 20 skills across 4 Pillars + Karpathy’s Vibe Coding Spectrum & openworker Earned Autonomy

Pillar 01: Building AI Apps

- **1.1 LLM Foundations:** Tokenization, context windows, cache hits, tool calling.
- **1.2 Grounding:** Prompt vs. on-demand retrieval ([chub](#)).
- **1.3 Agentic Systems:** Harnesses, MCP, memory, guardrails.
- **1.4–1.6 Evals & Prod:** Deterministic + LLM-judge evals.

Pillar 02: SE Fundamentals

- **2.1–2.2 Full-Stack & Data:** Schemas, transactions, state.
- **2.3 System Architecture:** Latency, consistency, modularity.
- **2.4 Security & Reliability:** Shift-left scanning, blast radius.
- **2.5 Prod Ops:** Observability & resilience steering.

Pillar 03: Coding Agents

- **3.1 Directing Workflow:** Upfront planning vs. fast execution.
- **3.2 Enabling Autonomy:** Permission tiers & subagents.
- **3.3 Reviewing Work:** Diffs, browser receipts, evals.
- **3.4–3.5 Customizing Harness:** [CLAUDE.md](#)/[GEMINI.md](#), Skills, Hooks.

Pillar 04: Shaping the Build

- **4.1 Driving the Build Loop:** Rapid Idea → MVP → Prod in small batches.
- **4.2 Product Decisions:** Scoping vertical slices.
- **4.3–4.4 High Agency:** “The Conductor Mindset”—orchestrating over abdicating.

Karpathy’s Vibe Coding Spectrum & andrewyng/openworker’s Ladder of Earned Autonomy

1. **Tier 0 – Pure Weekend Vibe Coding** ([MenuGen](#) / [llm-council](#) **hack**): Accept all diffs, ignore code → *Only for throwaway prototypes*.
2. **Tier 1 – Approval-Gated Steering** ([agy request-review](#) / [claude default](#)): Human reviews every edit and shell command.
3. **Tier 2 – Sandboxed Auto-Edit** ([agy proceed-in-sandbox](#) / [claude acceptEdits](#)): Auto-approves local file edits; gates network/shell.
4. **Tier 3 – Bounded Autonomous Loop** ([agy always-proceed](#) / [claude auto](#)): Reviewer-gated with **Hard Floors** (never self-approves destructive ops).

Lab 01: Architecture-Guided Steering & Earned Autonomy (prompt_steerer.py)

Walking through labs/lab_01_vibe_coding_and_earned_autonomy/ across agy and claude

Step 1 & 2: SE Prompt Scorer & Autonomy Tier Router

- **Auditing Naive vs. SE-Steered Prompts (score_prompt_steering):**

1. Flags underspecified “pure vibe” prompts lacking architectural constraints.
2. Scores explicit **Pillar 02 SE dimensions**: `latency_sla`, `consistency_model`, `failure_handling`, `security_boundary`, and `verification_command`.

- **Dual-Harness Permission Mapping (resolve_autonomy_tier):**

Maps risk profiles to `~/.gemini/antigravity-cli/settings.json` (`request-review`, `proceed-in-sandbox`, `always-proceed`, `strict`) and `.claude/settings.json` (`default`, `acceptEdits`, `plan`, `auto`).

Step 3 & 4: Hard Floors & Circuit-Breaker Verification

- **openworker Hard Floor Invariant:** Even in `always-proceed` or `auto` mode, destructive commands (`rm -rf`, `git push -force`, prod DB drops, secret exfiltration) hit an unbypassable `HARD_FLOOR_BLOCK`.

- **Earned Autonomy Graduation:** Repeated verified runs promote commands from one-off approval → standing allowlist rule; consecutive reviewer denials trip the **circuit breaker** back to `plan` / `request-review`.

- **Run Self-Diagnosis in < 0.05s:**

```
./labs/lab_01_vibe_coding_and_earned_autonomy/self_diagnose.sh
```

Lab 01 Key Takeaway: Autonomy Is Earned by Harness Boundaries, Never Granted Blindly

- Combine **Pillar 02 Software Engineering Vocabulary** in your prompt with **Pillar 03 Permission Tiers & Hard Floors** in your harness settings before letting an agent execute multi-file changes.

Anatomy of a Coding Agent on the Wire (aisuite & Stanford CS146S Lecture 2)

The 3-Layer Agent Stack, 200-line minimal agent vs. production JSONL wire trace, and “Swiss Army Knife” primitives


The 3-Layer Agent Stack (andrewyng/aisuite) & Wire Trace

- Layer 1:** Chat Completions API (Stateless `messages[]` → `provider:model`).
- Layer 2:** Agents API · Toolkits (`files`, `git`, `shell`) · MCP.
- Layer 3:** Agent Harness (Context compiler, permissions, hooks, compaction).
- 200-Line Minimal Agent vs. Production HTTP/JSONL Payload:**
- *Minimal Loop:* `while tool_calls: run(read|write|edit|bash).`
 - *Prod Wire Trace:* `System Prompt` + `CLAUDE.md/GEMINI.md` + `Rules` + `Skill YAML Catalog` + `Native Tool Schemas` + `mcp__* Schemas` + `Message History`.

Why “Swiss Army Knife” Primitives Beat 150 CRUD Tools

- CS146S Lecture 2 Wire-Trace Discovery:**
- Mounting a naive 150-endpoint GitHub/Jira MCP server injects **10,000+ lines of JSON schema definitions** for just **600 lines of conversation**!
- The “Swiss Army Knife” Consolidation Principle:**
- Production harnesses rely on **5–8 composable primitives**: `Bash`, `Read`, `Write`, `Edit`, `Grep`, `Glob`, and `LSP`.
 - Frontier models have deep post-trained fluency in `git`, `gh`, `jq`, `ripgrep`, and `pytest` via `Bash`—saving > 85% of schema tokens!

Wire-Trace Token Breakdown: Bloated 150-Tool Harness vs. Lean “Swiss Army Knife” Harness



Google Antigravity (agy) 101: Mission Control, IDE, Artifacts & CLI

Grounded in `getting-started-google-antigravity`, `antigravity-cli-hands-on` & `build-and-deploy-gcp-with-antigravity`

1. Mission Control & IDE

Antigravity 2.0 Desktop & Editor:

- **Agent Manager (Mission Control):** Spawn parallel workspace agents across multi-folder Projects.
- **In-Editor Shortcuts:** `Cmd+L` (AI chat pane), `Tab to Jump` (next logical edit point), and inline `Explain and fix`.
- **/browser Chrome Subagent:** Live DOM inspection, console log capture, and visual verification screenshots.

2. First-Class Artifacts

Verifiable Deliverables over Raw Chat:

- **Task List:** Live structured progress checklist updated autonomously.
- **Implementation Plan:** Blueprint supporting **Google Docs-style inline comments** before clicking **Proceed**.
- **Walkthrough (Run Receipt):** Post-implementation proof bundling diffs, test logs, and browser recordings.

3. Antigravity CLI (agy)

Terminal & Headless Automation:

- **Core CLI:** `agy -version`, `agy models`, `agy -p "..."` `-model ...`, and `!` Shell Mode toggle.
- **Settings:** `~/.gemini/antigravity-cli/settings.json` & `/config`.
- **Power Slash Commands:** `/plan`, `/artifact`, `/grill-me`, `/goal`, `/schedule`, `/stats`, `/memory`, `/skills`, `/compress`, `/clear`.

Antigravity Customization Hierarchy (GEMINI.md + AGENTS.md + .agents/)

- **Workspace Configs:** Root `GEMINI.md` & `AGENTS.md` → `.agents/rules/*.md` → `.agents/skills/<name>/SKILL.md` → `.agents/agents/*.md` → `.agents/workflows/*.md`.
- **Persistent Knowledge Items:** Stored at `~/.gemini/antigravity/knowledge/<project>/KNOWLEDGE.md` and managed via `/memory show|add|refresh`.

Anthropic Claude Code (claude) 101: Surfaces, Modes & Slash Commands

Grounded in Anthropic Academy's `claude-code-101` and DeepLearning.AI's *Claude Code: A Highly Agentic Coding Assistant*

1. Core Loop & Surfaces

Gather Context → Act → Verify:

- Runs natively in **Terminal CLI** (`claude`), **VS Code / JetBrains, Desktop App**, and **Web**.
- **Context Pinning:** Use `@file` and `@dir/` to inject exact source files without noisy grep searches.
- **Direct Bash Mode:** Prefix any command with `!` (e.g., `!pytest -q`) to run shell commands directly in transcript.

2. 3-Stage In-Editor Cycle

Interactive Steering & `Shift+Tab`:

- **`Shift+Tab` Mode Cycling:** Toggle instantly between `Normal` → `Auto-Accept Edits` → `Plan Mode`.
- **Real-Time Interruption:** Queue follow-up messages mid-flight or press `Esc` to halt a runaway tool loop.
- **Headless CI/Script Mode:** Run `claude -p "review diff" -output-format json`.

3. Context & Time-Travel

Essential Session Commands:

- `/init` & `/memory`: Bootstrap and edit `CLAUDE.md` and project `MEMORY.md`.
- `/context`: Inspect live token usage across system, files, MCP, and messages.
- `/compact [focus]` & `/clear`: Compress or wipe history between tasks.
- `/rewind (Esc+Esc)` & `/btw`: Checkpoint rollback and zero-history side queries.

1-to-1 Command Translation Cheat Sheet (`agy` ↔ `claude`)

- **Inspect Context:** `agy /stats` ↔ `claude /context`
- **Compress / Reset:** `agy /compress, /clear` ↔ `claude /compact, /clear`
- **Plan Mode:** `agy /plan` (Implementation Plan Artifact) ↔ `claude Shift+Tab` (`Ctrl+G` editor)
- **Shell Toggle:** `!` in both CLIs.

Lab 02: Wireshark-Style Harness Trace Inspector (`wire_trace_inspector.py`)

Inspecting JSONL wire traces, tool schema bloat, and loop circuit breakers in `labs/lab_02_agent_loop_and_harness_wire_trace/`

Step 1 & 2: Parsing the Wire Trace & Schema Overhead

- **Inspecting Full HTTP/JSONL Payloads (`inspect_wire_trace`):**

1. Dissects a captured agent request into `system_prompt`, `standing_context` (`CLAUDE.md/GEMINI.md`), `native_tools`, `mcp_tools`, and `messages`.
2. Computes exact token shares and flags when `mcp_tools` schema bloat exceeds the **15% schema budget cap**.

- **Consolidating to Swiss Army Knife Primitives (`prune_to_core_primitives`):**

Replaces redundant REST CRUD tool definitions with composable `Bash/Read/Edit/Grep` primitives, recovering $> 80\%$ of fixed startup overhead.

Step 3 & 4: Runaway Loop Detection & Self-Diagnosis

- **Detecting Spin Loops on the Wire (`detect_runaway_tool_loop`):**

1. Identifies repeated identical tool invocations (`tool_name + canonical_args`) or oscillating edit-fail-revert cycles across consecutive turns.
2. Trips a deterministic `CircuitBreakerError` after $N = 3$ repeated failures and emits a structured prompt reset recommendation.

- **Verify Lab 02 in $< 0.05\text{s}$:**

```
./labs/lab_02_agent_loop_and_harness_wire_trace/self_diagnose.sh
```

Lab 02 Key Takeaway: Never Debug an Agent Blind—Inspect the Wire Payload

- Whenever an agent behaves erratically or runs slow, inspect the wire trace (`/stats` or `/context`): 90% of failures stem from hidden schema bloat or polluted message turns.

Task Complexity vs. Workflow Rigor & Defeating “Silent Assumption Syndrome”

Calibrating process overhead to task blast radius (Stanford CS146S Lecture 3) and Socratic Inversion-of-Control via `/grill-me`

Task Complexity vs. Workflow Rigor Spectrum

- **Level 1 – Copy / Typo / Single-Function Tweak:**

Workflow: Direct vanilla prompt with `@file` pin (~1 turn).

- **Level 2 – Small Multi-File Bugfix or Endpoint (2–4 files):**

Workflow: Single **Plan Mode** pass (`agy /plan` or `claude Shift+Tab`) → Execute → Test.

- **Level 3 – Medium Feature or Schema Change (4–12 files):**

Workflow: Full 5-Step **RePPIT Loop** (Research → Propose → Plan → Implement → Verify).

- **Level 4 – Large Cross-Module Refactor / Migration (>12 files):**

Workflow: Multi-session stacked vertical slices with isolated `git worktree` subagents.

Defeating “Silent Assumption Syndrome” via `/grill-me`

- **Why RLHF Models Guess Silently:** Standard LLMs are post-trained to be “helpful immediately,” causing them to silently guess auth schemes, pagination limits, error formats, and retry policies.

- **Socratic Inversion of Control (`/grill-me` & Plan Mode):**

Instead of asking the agent to write code immediately, invert control:

`"Interview me relentlessly about every ambiguity, edge case, and architectural trade-off in this feature. Do NOT write code until all questions are resolved."`

- Surfaces hidden requirements in 2 minutes before a single line of code is written.

Skills Map 3.1 Rule: Match Upfront Planning Depth to Blast Radius

- Under-planning a medium feature causes architectural rewrite loops; over-planning a 3-line CSS fix wastes tokens. Use the 4-level spectrum above on every task.

The 5-Step RePPIT Loop (themodernsoftware.dev CS146S Lecture 3)

Research → Propose (+ Context Reset!) → Plan (+ Out of Scope) → Implement (Model Routing) → Verify & Reflect

1. RESEARCH

Objective Discovery:

- Document *only what exists today* with exact `file:line` citations.
- Strictly **forbidden from proposing solutions** during research.
- Builds both compressed agent context and the **human's mental model**.

2. PROPOSE + RESET

Orthogonal Options:

- Force **2 orthogonal proposals** (e.g., sync DB vs. async queue) with trade-offs.
- **Post-Proposal Context Reset:** Once Proposal B is picked, **purge Proposal A from context** so rejected tokens never confuse coding!

3. PLAN

Locked Blueprint:

- File-and-line edit list in `Implementation Plan` Artifact or `Ctrl+G`.
- Mandatory **Out of Scope / Do NOT Touch** guardrail section blocking scope creep.
- Explicit test verification matrix.

4. IMPLEMENT

Phase Model Routing:

- Use a **frontier reasoning model** for Steps 1–3 (*Research, Propose, Plan*).
- Route Step 4 (*Implement*) to a **fast, cost-efficient coding model** once the design is locked!

5. VERIFY & REFLECT

Critique & Evidence:

- Run `translation-agent Reflection pass` (`Draft → Critique → Refine`).
- Execute unit tests + live `/browser` / Playwright visual evidence in `Walkthrough`.

Two High-Leverage RePPIT Rules That Eliminate 80% of Agent Hallucinations

- **Rule 1 – Post-Proposal Context Reset:** If an agent generates 1,500 tokens describing Rejected Proposal A and 1,500 tokens describing Chosen Proposal B, leaving Proposal A in the chat history causes the model to accidentally blend Proposal A's types into Proposal B's implementation. Always compact/reset after choosing!
- **Rule 2 – Explicit “Out of Scope / Do NOT Touch” Section:** Coding agents love “helpful” drive-by refactors. Listing forbidden files/modules in the plan keeps diffs surgical.

Side-by-Side Daily Execution in agy vs. claude & Time-Travel Recovery

Reviewing plans inline, capturing Run Receipts, and recovering instantly from bad trajectories

Google Antigravity (agy) Execution & Recovery

- **Inline Plan Co-Editing:** Open the generated `Implementation Plan` Artifact, highlight any over-engineered function or missing edge case, and leave a **Google Docs-style inline comment** before clicking **Proceed**.
- **Browser Verification Receipt:** Invoke `/browser` to upload synthetic test data, inspect console errors, and attach screenshots automatically to the `Walkthrough` Artifact.
- **Time-Travel Checkpoint Rollback:** Click `Undo changes up to this point` on any earlier step in the trajectory to revert both workspace files and conversation state simultaneously.

Anthropic Claude Code (claude) Execution & Recovery

- **Interactive Plan Editing (`Ctrl+G`):** Enter Plan Mode via `Shift+Tab`, let Claude draft the plan, and press `Ctrl+G` to open the plan directly in your text editor to trim scope before execution.
- **Playwright MCP Verification:** Connect Playwright MCP to launch headless browser checks and verify rendered DOM state alongside `!pytest`.
- **Time-Travel via `/rewind` (`Esc+Esc`):** Press `Esc` twice (or run `/rewind`) to open the checkpoint menu and restore code only, conversation only, or both.

Why Time-Travel Beats “Arguing in Chat”

- When an agent takes a wrong architectural turn, **never spend 5 turns arguing with it in chat** (which fills the context window with broken code and error logs). Rewind to the plan checkpoint, tighten the constraint in 1 line, and re-run cleanly!

Lab 03: End-to-End RePPIT Orchestrator (`reppit_orchestrator.py`)

Walking through `labs/lab_03_reppit_workflow_and_reflection/` step by step

Step 1, 2 & 3: Research Audit, Proposals & Context Reset

- **Step 1 – Pure Research Validator (`validate_research_artifact`):**

Verifies every claim cites an existing `file:line` and rejects premature code snippets or solution proposals.

- **Step 2 – Orthogonal Proposal & Context Reset (`select_proposal_and_reset`):**

Requires ≥ 2 distinct architectural proposals with explicit trade-off matrices, locks the chosen proposal, and **purges rejected proposal tokens** from the active context state.

- **Step 3 – Plan Gate (`validate_implementation_plan`):**

Enforces non-empty `target_files`, `verification_steps`, and `out_of_scope` guardrails.

Step 4 & 5: Model Routing, Reflection & Run Receipt

- **Step 4 – Phase-Based Model Router (`route_phase_model`):**

Routes `RESEARCH`, `PROPOSE`, and `PLAN` to the frontier reasoning tier, and routes `IMPLEMENT` to the fast execution tier while checking that zero `out_of_scope` files were modified.

- **Step 5 – `translation-agent` Reflection Pass (`run_reflection_pass`):**

Executes `Draft` $\rightarrow$ `Critique` $\rightarrow$ `Refine` against domain invariants and compiles the verified `Walkthrough` Run Receipt.

- **Verify Lab 03:** `./labs/lab_03_reppit_workflow_and_reflection/self_diagnose.sh`


Lab 03 Key Takeaway: Discipline at the Proposal & Plan Gates Makes Implementation Trivial

- By purging rejected proposals, locking `Out of Scope` files, and running a reflection pass before commit, Lab 03 eliminates scope creep and silent regressions.

Context Window Physics (`rmvDxxNubIg`): The 40% “Dumb Zone”

Stateless NextAction = LLM(ContextWindow) physics and why “not all 1M tokens are created equal”

Reasoning & Instruction-Following Quality vs. Context Utilization



Dex Horthy's First Law of Coding Agents (`rmvDxxNubIg`)

- **Stateless Function Physics:** Every single turn, an LLM computes NextAction = LLM(ContextWindow) from scratch. The only lever you have to improve the next action is the **signal-to-noise ratio of the context window**.

- **What Floods Context Past 40% into the “Dumb Zone”:**

1. Unfiltered `grep` or `find` dumps across `node_modules/` or build logs.
2. Verbose test/training stdout (5,000 lines of progress bars and tracebacks).
3. Multi-turn “That’s still wrong, try again” correction loops.
4. Bloated MCP tool schemas and monolithic 500-line instruction files.

4-Dimension Context Optimization Hierarchy: Correctness > Completeness > Size (<40%) > Trajectory

- 1. No Incorrect Information (worst poison) → 2. No Missing Information → 3. Minimal Noise (<40% capacity) → 4. Clean Forward Trajectory (never argue with a polluted history).

Context Backpressure (karpathy/autoresearch & rendergit) & Side-Queries

Preventing stdout floods via shell redirection, dual-view repo packing, and zero-history `/btw` queries

1. Output-Redirection Backpressure

Karpathy's `autoresearch` Rule:

- Never let a test suite or training job stream raw output into the LLM context window (*"do NOT use tee or let output flood your context"*).

- **Canonical Pattern:**

```
uv run train.py > run.log 2>&1
```

```
grep "^val_bpb:|^peak_vram:" run.log
```

- On crash, read only `tail -n 50 run.log`. Compresses 50,000 log tokens to **40 tokens!**

2. karpathy/rendergit Dual View

Human HTML vs. LLM CXML View:

- When ingesting a small repository or module, `rendergit` provides two tailored representations:
 1. **Human View:** Syntax-highlighted static HTML with navigable directory tree sidebar.
 2. **LLM CXML View:** Clean, token-counted XML document boundaries (`<document path="...">`) stripping binary/lockfile noise.
- Pair with `karpathy/nanochat` minimalism.

3. Sandboxed Side-Queries (`/btw`)

Zero-Footprint Tangents:

- Mid-implementation, you often need to ask a quick question (*"What flag does `pytest` use for Doctests?"*).
- Asking in the main chat appends the tangent to every future turn.
- **Claude Code `/btw <question>`:** Answers in a modal using current context, then **discards the Q&A from history!**
- Monitor budget via `/context` & `/stats`.


Wrapper Script Pattern for Deterministic Backpressure

- Wrap noisy build/test commands in a helper script that outputs `[PASS] 24 tests (0.12s)` on success and prints *only* the failing assertion traceback on error.


Frequent Intentional Compaction (FIC), RPI & The Human Leverage Pyramid

Sawtooth context management via `research.md` / `plan.md` handoffs and why 1 line of plan review saves 1,000 lines of code

Sawtooth Context Utilization: Naive Accumulation vs. FIC (RPI)



The Human Attention Leverage Pyramid (rmvDxxNubIg)



How Frequent Intentional Compaction (FIC) Works Between RPI Phases

- At the end of **Research**, write verified findings + exact `file:line` citations to `research.md` and run `/clear`.
- Start **Plan** with a fresh window reading only `research.md` → write `plan.md` → run `/clear` → execute **Implement** in the Smart Zone (< 25% utilization)!

Lab 04: Context Budget Profiler & RPI Compaction Engine (`context_compactor.py`)

Walking through `labs/lab_04_context_physics_and_rpi_compaction/` step by step

Step 1 & 2: 40% Dumb-Zone Profiler & Log Backpressure

- **Step 1 – Context Budget Profiler** (`profile_context_window`):

1. Computes total utilization ratio across system instructions, tool schemas, file reads, and tool outputs against the model's context window.
2. Flags `DUMB_ZONE_WARNING` the moment utilization crosses **0.40 (40%)**.

- **Step 2 – Output-Redirection Backpressure** (`apply_log_backpressure`):

Filters multi-thousand-line test/training stdout logs down to summary metrics (`val_bpb`, `passed/failed` counts) + the final 30 traceback lines on failure.

Step 3 & 4: RPI Artifact Compaction & Self-Diagnosis

- **Step 3 – Intentional Compaction Engine** (`compact_to_rpi_artifact`):

1. Distills a noisy exploration transcript into a structured `research.md` or `plan.md` handoff artifact.
2. Validates that every retained claim preserves its canonical `path/to/file.py:L10-L25` citation and strips failed search dead-ends.
3. Resets active session utilization from **>65% back down to <12%**.

- **Verify Lab 04 in < 0.05s:**

```
./labs/lab_04_context_physics_and_rpi_compaction/self_diagnose.sh
```

Phase 1 Complete! You Now Command Daily Vibe Coding & Context Hygiene in `agy` and `claude`

- Next, in **Phase 2 (Modules 5–8)**, we transition from operating the agent interactively to **Harness Engineering**—designing the rules, skills, MCP code mode, subagent firewalls, and ratchets around the model.

Entering Phase 2: Harness Engineering (Agent = Model + Harness)

Birgitta Böckeler & Kief Morris’s 2x2 Cybernetic Control Matrix (martinfowler.com) and Why Loop vs. How Loop

The 2x2 Cybernetic Control Matrix (Böckeler & Morris, martinfowler.com)

	GUIDES (Feedforward / Before Action)	SENSORS (Feedback / After Action)
COMPUTATIONAL (Deterministic CPU)	1. Deterministic Guides - OS/Container Sandboxes (OpenShell) - Tool allowlists & PreToolUse Exit-2 gates - Codemods & static project templates	2. Deterministic Sensors - Compiler & Static Typecheckers - Unit/Integration Tests & PostToolUse - Remediation-Aware Structural Linters
INFERENTIAL (Semantic LLM)	3. Inferential Guides CLAUDE.md / GEMINI.md (60-line map) - Glob-scoped rules/*.md & SKILL.md - Curated API Grounding (chub)	4. Inferential Sensors - Multi-Agent PR Review Subagents - Karpathy’s llm-council Peer Review translation-agent Reflection Pass

Kief Morris: “Why Loop” vs. “How Loop”

- Human-Owned “Why Loop”:** Defines business outcomes, domain architecture, program design (types/interfaces), and trade-offs.
- Agent-Executed “How Loop”:** Explores files, drafts code, runs tests, and self-repairs inside the harness.
- The “On-the-Loop” Harness Flywheel:**
Whenever a defect escapes the How Loop, **do not merely patch the code line**—upgrade one of the 4 Cybernetic Quadrants so that entire class of failure can never recur!


The Golden Rule of Harness Engineering: Prefer Computational (CPU) Checks Over Inferential (LLM) Hopes

- Never write a 20-line prose paragraph in [CLAUDE.md](#) begging the model to format imports or forbid a deprecated API if a 10-line **Deterministic Linter** or [PreToolUse](#) **Hook** can enforce it with 100% certainty at zero token cost!

The Three Context Primitives, The 150-Instruction Cliff & The 60-Line Map Rule

Prompts vs. Rules vs. Skills, why 500-line instruction files overfit (OpenAI & CS146S), and progressive `@docs/` imports

The 150-Instruction Compliance Cliff (Why Monolithic Rules Fail)



The Three Context Primitives:

1. **Prompts / Slash Workflows:** Ephemeral per-turn user instructions (`/plan`, `/grill-me`).
2. **Rules** (`CLAUDE.md/GEMINI.md`, `rules/*.md`): Standing constraints (always-on or glob-scoped).
3. **Skills** (`SKILL.md`): On-demand procedural manuals (`~100 tok` YAML until triggered).

The Minimalist 60-Line Map Rule (OpenAI & CS146S)

- **Why 500-Line Files Hurt:** Over-stuffed instruction files consume instruction budget, contradict stale modules, and **overfit agent exploration**.
- **What Belongs in Root** `CLAUDE.md` / `GEMINI.md` (≤ 60 lines):
 1. **Exact Build / Test / Lint / Server Commands** (e.g., `./self_diagnose_all.sh`).
 2. **Concise Architecture & Directory Map** (top-level package boundaries).
 3. **Progressive `@docs/*.md` Pointers** (loaded only when relevant to the task).
 4. **Cross-Tool Symlink:** Symlink `CLAUDE.md` $\leftrightarrow$ `AGENTS.md` (and `GEMINI.md`).
- **Path-Scoped Rules:** Move folder-specific rules into `.claude/rules/api.md` (`paths: ["src/api/**"]`) and `.agents/rules/api.md`.

Context Economy Rule: Treat Root Instruction Lines Like L1 CPU Cache

- If a rule applies to $< 20\%$ of tasks (e.g., database migration conventions or frontend CSS tokens), move it out of root `CLAUDE.md/GEMINI.md` into a path-scoped `rules/*.md` file or an on-demand `SKILL.md`.

Cross-Session Auto-Memory & Retrospectives (MEMORY.md vs. KNOWLEDGE.md)

Persisting durable architectural decisions across sessions while pruning stale assumptions

Claude Code Auto-Memory (MEMORY.md)

- **Storage Path:** `~/.claude/projects/<git-repo>/memory/MEMORY.md` (shared across all `git worktree` checkouts of the repository).
- **Startup Cap:** Only the first **200 lines / 25 KB** of `MEMORY.md` are auto-loaded; detailed sub-notes route to `topics/<slug>.md`.
- **Interactive Inspection:** Run `/memory` to inspect, edit, or delete entries when project architecture changes.

Google Antigravity Knowledge Items (KNOWLEDGE.md)

- **Storage Path:** `~/.gemini/antigravity/knowledge/<project>/KNOWLEDGE.md` with structured topic summaries and artifact pointers.
- **Interactive CLI Commands:** Manage standing knowledge directly via `/memory show`, `/memory add`, and `/memory refresh`.
- **Post-Run Retrospectives (Skills Map 3.4):** After non-trivial sessions, capture discovered gotchas and **update downstream entries** when assumptions change.

Preventing Memory Poisoning: Provenance Tagging & Stale-Assumption Pruning

- Never let an agent record speculative hypotheses or transient debugging state as permanent rules. Require explicit **source provenance** (user instruction vs. commit-verified fact) and prune superseded entries on architecture updates.

Lab 05: Context Compiler & Persistent Memory Architecture (`context_compiler.py`)

Walking through `labs/lab_05_context_primitives_and_memory/` step by step

Step 1 & 2: 60-Line Map Compiler & Path-Scoped Rules

- **Step 1 – Monolith-to-Map Refactorer (`compile_root_context_map`):**
 1. Audits a bloated 400-line instruction file and enforces the strict **60-line root cap**.
 2. Extracts domain-specific sections into `@docs/<topic>.md` lazy imports with path-traversal security checks (`relative_to(project_root)`).
- **Step 2 – Glob-Scoped Rule Matcher (`match_scoped_rules`):**

Loads `.claude/rules/*.md` and `.agents/rules/*.md` only when the active file set matches the rule's frontmatter `paths: globs`.

Step 3 & 4: Worktree Memory Store & Self-Diagnosis

- **Step 3 – Cross-Session Memory & Retrospectives (`MemoryStore`):**
 1. Resolves canonical git-worktree memory paths for both `MEMORY.md` and `KNOWLEDGE.md`.
 2. Enforces the **200-line / 25 KB index cap** (routing overflow into `topics/*.md`) and superseding stale assumptions during post-run retrospectives.
- **Verify Lab 05 in < 0.05s:**

```
./labs/lab_05_context_primitives_and_memory/self_diagnose.sh
```

Lab 05 Key Takeaway: 82% Startup Token Reduction with Zero Loss of Domain Guidance

- Refactoring a monolithic prompt into a 60-line root map + glob-scoped rules + topic-overflow memory keeps baseline context under 4% of the window.

Skills as On-Demand Procedural Manuals (agentskills.io)

3-Level Progressive Disclosure (YAML → SKILL.md → scripts/references) and the 3-Stage Skill Maturity Lifecycle

Level 1: Discovery Catalog

YAML Frontmatter (~100 Tokens):

- Only `name` (lowercase-hyphenated matching folder) and `description` load at session startup.
- Must state both positive triggers (`Use when...`) and negative boundaries (`Don't use for...`).
- 20 installed skills cost just ~2,000 startup tokens!

Level 2: Core Manual

SKILL.md Body (≤ 500 Lines):

- Loaded on-demand only when invoked via `/skill-name` or matched by the semantic router.
- Contains concrete workflow steps, exact CLI recipes, and **experience-grounded gotchas**.
- Links one level deep to Level 3 files.

Level 3: Bundled Resources

scripts/, references/, assets/:

- **scripts/**: Deterministic Python/Bash verifiers executed via shell without reading code into context!
- **references/**: Deep domain tables/schemas read selectively per sub-task.
- Costs **0 tokens** until explicitly accessed.

The 3-Stage Skill Maturity Lifecycle (Skills Map 3.4): Prototype → Team Preview → Production GA (+ Pruning)

- **1. Personal Prototype** (captured from a real debugging trajectory) → **2. Team Preview** (checked into `.agents/skills/` & `.claude/skills/`) → **3. Production GA** (gated by paired `benchmark.json` evals) → **4. Retire/Prune** when newer frontier model weights absorb the capability natively!

Curated API Grounding & Self-Improving Annotations (andrewyng/context-hub)

How @aisuite/chub (get-api-docs/SKILL.md) eliminates stale-API hallucinations and compounds session discoveries

1. Stale Pretraining Problem

- LLMs hallucinate deprecated SDK methods (e.g., old `openai.ChatCompletion.create` or stale Stripe webhook signatures).
- Raw web search pulls noisy HTML ads and mismatched versions.

2. chub Curated Fetch

- Agent triggers `get-api-docs/SKILL.md` and runs:

```
chub search stripe  
chub get stripe/api -lang py
```
- Injects clean, versioned, language-specific Markdown docs on demand.

3. Local Self-Annotation

- When the agent discovers a non-obvious SDK quirk during testing, it persists a local note:

```
chub annotate stripe/api "Requires  
raw bytes payload for HMAC check"
```
- Stored locally for future sessions!

4. Guarded Re-Use & Feedback

- Future sessions run `chub get ... -with-annotations` (wrapped as untrusted context so notes cannot inject commands).
- Emits quality signal upstream via `chub feedback up|down`.

Why andrewyng/context-hub Outperforms Static RAG and Raw Documentation Scraping

- By combining **curated version-pinned Markdown packages** with **local agent-authored annotations** (`chub annotate`), the harness turns every hard-won API debugging discovery into permanent institutional memory across sessions.


Ergonomic MCP Tool Design & Meta-MCP “Code Mode” (>90% Token Savings)

Why 1-to-1 REST CRUD wrappers fail and how Stanford CS146S Lecture 4’s 4 Rules cut tool tokens by >90%

The 4 Rules of Ergonomic Agent Tools (CS146S Lecture 4)

- 1. Design Around Outcomes, Not Granular REST CRUD:**
Replace 3 chained REST lookups (`get_user` → `get_wallet` → `post_tx`) with 1 outcome tool (`send_payment_to_email`) or consolidated `search_resource(type=...)`.
- 2. Enforce Strict Typed Schemas (`Pydantic` / `TypeScript`):**
Ban `*args`, `**kwargs` and untyped `dict` payloads; constrain enums and value bounds.
- 3. Return Actionable, Self-Repairing Error Messages:**
Instead of `HTTP 400 Bad Request`, return exact parameter remediation instructions.
- 4. Expose Meta-MCP “Code Mode” (`search` + `execute`):**
Instead of loading 150 tool schemas upfront, expose only `search(query)` (to fetch 1–2 relevant schemas on demand) and `execute(script)` (running chained calls in a sandbox VM).

Context Tokens Consumed by 5-Step Multi-Tool Workflow




Why Meta-MCP “Code Mode” Is the Architecture Standard for Large Tool Catalogs

- LLMs are far better at writing a 6-line Python/TypeScript script inside `execute()` that loops over paginated API responses than issuing 12 sequential JSON tool-call round-trips across the context window.

Empirical Skill Science (SkillsBench arXiv:2602.12670) & eval-viewer

Why focused curated skills add **+16.6pp** (2–3 skills peak at **+19.0pp**) while self-generated skills hurt (–8.1 to –11.5pp)

SkillsBench Pass-Rate Delta vs. No-Skill Baseline (87 Tasks, 9,396 Runs)



Why Self-Generated Skills Fail & How to Evaluate

- **Why Self-Generated Skills Hurt (–8.1 to –11.5pp):** When asked to write its own skill before solving a task, an LLM emits generic pretraining platitudes or locks in hallucinated API flags.
- **Why > 10 Skills Degrade Routing:** Overlapping descriptions cause trigger collisions and blow past the active context budget.
- **Paired eval-viewer Ablation (benchmark.json):**
 1. Run every eval prompt **paired:** `with_skill` vs. `without_skill` baseline.
 2. Measure `delta_pass_rate`, `delta_tokens`, and trigger discrimination (`should_trigger` vs. `should_not_trigger` false-alarm rate).

Experience Before Theory: Craft Skills from Real Failure Trajectories & Keep Active Skills at 2–3

- Only promote a skill to the team repository if its paired `benchmark.json` demonstrates a strictly positive `delta_pass_rate` without inflating token consumption on non-matching tasks.

Lab 06: Progressive Skills, chub Annotations & Meta-MCP (`skill_and_mcp_harness.py`)

Walking through `labs/lab_06_skills_chub_and_meta_mcp_code_mode/` step by step

Step 1 & 2: `agentskills.io` Validator & chub Grounding

- **Step 1 – 3-Level Skill Package Validator (`validate_skill_package`):**

Verifies YAML frontmatter (`name`, positive/negative `description` triggers), $\leq$ 500-line `SKILL.md` body, one-level-deep `references/` links, and rejects self-generated filler platitudes.

- **Step 2 – Self-Annotating Doc Store (`ChubDocRegistry`):**

Implements `search()`, versioned `get(id, lang, with_annotations=True)`, local `annotate(id, note)` persistence, and `feedback(id, vote)`.

Step 3 & 4: Meta-MCP “Code Mode” & `benchmark.json`

- **Step 3 – Meta-MCP `search` + `execute` Harness (`MetaMcpCodeMode`):**

1. Exposes only 2 meta-tools (`search_tools` and `execute_code_plan`) over a catalog of 120 backend tools.
2. Proves >90% **schema token reduction** while enforcing strict typed input validation and actionable remediation errors.

- **Step 4 – Paired Ablation Report & Self-Diagnosis:**

`./labs/lab_06_skills_chub_and_meta_mcp_code_mode/self_diagnose.sh`

Lab 06 Key Takeaway: On-Demand Progressive Disclosure Applies to Both Skills and MCP Tools

- Loading skills via Level-1 YAML and loading tools via Meta-MCP `search` + `execute` ensures your agent scales to 100+ capabilities while keeping > 85% of its context window free.

The Interaction Method Decision Matrix & Subagents as Context Firewalls

Choosing between Conversational Prompts, One-Off Prompts, and Subagents (rmvDxxNubIg & Skills Map 3.2)

1. Conversational Prompts

When to Use Interactive Turns:

- High-ambiguity requirement discovery (`/grill-me`) and architectural trade-off selection.
- Iterative UI/UX refinement where human judgment steers each step.
- *Warning:* Never use conversational turns for broad 30-file search sweeps!

2. One-Off Bounded Prompts

When to Use Single-Shot / Headless:

- Well-specified transformations with deterministic verification (`agy -p` / `claude -p`).
- Generating commit messages, changelog entries, or single-file unit test suites.
- Zero multi-turn history accumulation.

3. Subagents = Context Firewalls

Dex Horthy's Subagent Law:

- **Never anthropomorphize subagents** into cute role-play personas (*"Frontend Bob"*, *"PM Alice"*).
- Subagents exist for one physical purpose: **Context Isolation**.
- Child burns **50,000 tokens** searching files/logs → returns a **500-token cited summary** to the parent window!

Context Firewall Compression Ratio: 50,000 Child Exploration Tokens → 500 Parent Handoff Tokens (100:1 Savings)

- Because the child subagent's context window is garbage-collected the moment it returns its structured findings, the parent orchestrator stays at < 15% context utilization!

Anthropic’s 7-Phase feature-dev & Stanford CS146S Multi-Agent PR Triage

Orchestrating code-explorer, code-architect, and 5 specialized code-reviewer subagents

Phases 1–3: Explore & Clarify

- **Phase 1 (Discovery):** Frame feature intent.
- **Phase 2 (code-explorer):** 2–3 parallel **read-only** sub-agents (`Read`, `Grep`, `Glob` only; `Write/Edit` forbidden) map execution paths with `file:line` citations.
- **Phase 3 (Clarification Gate):** Hard STOP until all open questions are answered.

Phases 4–5: Architect & Build

- **Phase 4 (code-architect):** Evaluates Minimal-Change vs. Clean-Architecture vs. Pragmatic lenses; locks file-and-line blueprint + `Out of Scope` list.
- **Phase 5 (Implementation):** Executes surgical edits against the approved blueprint and runs deterministic test suites.

Phases 6–7: 5-Agent PR Triage

- **Phase 6 (code-reviewer Fan-Out):** Spawns 5 parallel review subagents: **Security**, **Performance**, **Architecture**, **Style**, and **Edge Cases**.
- Filters out noise below **confidence** ≥ 80 .
- **Phase 7 (Summary):** Synthesizes final Run Receipt.

Stanford CS146S Multi-Agent PR Triage Taxonomy (Confidence ≥ 80 Gate)

- **[RED] Must Fix (Confidence ≥ 90 , Critical/High):** Security vulnerabilities, data-loss bugs, race conditions, or broken invariants (blocks merge).
- **[YELLOW] Recommended (Confidence ≥ 80 , Medium):** Missing boundary test cases, N+1 query inefficiencies, or error-handling gaps.
- **[GREEN] Consider (Confidence ≥ 80 , Low):** Optional readability simplifications; everything with confidence < 80 is automatically suppressed!

Karpathy's 11m-council & Agentic VCS: git worktree vs. Jujutsu (jj)

3-stage anonymized peer ranking + git worktree (Provisioning Cost) vs. colocated jj (Integration Cost, wavect.io)

Karpathy's 11m-council: 3-Stage Verification Protocol

- **Why Single-Model Self-Review Fails:** A single model evaluating its own design exhibits self-preference bias and misses its own blind spots.
- **Stage 1 – Independent First Opinions:** Query K diverse models in parallel.
- **Stage 2 – Anonymized Peer Review:** Strip model identities (`Response A..C`) so reviewers score accuracy and simplicity without playing favorites.
- **Stage 3 – Chairman Synthesis:** Aggregate Borda count rankings + peer critiques into one unified architectural decision.

Agentic VCS: git worktree vs. Jujutsu (jj) (wavect.io)

- **git worktree (Cuts Provisioning Cost):** Shares one `.git` object store across Antigravity Mission Control & `claude --worktree` (+ `--filter=blob:none`).
- **Colocated Jujutsu (`jj git init --colocate` + `jj workspace add`):**
 1. **Working Copy = Commit (@) & Stable Change IDs:** Zero-ceremony auto-snapshots on every command; Change IDs survive rewrites.
 2. **First-Class Conflict Algebra (ΔM):** Rebases **never abort mid-loop**; conflicts store in commits, descendants auto-evolve, and N-parent stacks (`jj new A B`) rebase cleanly.
 3. **Transactional Time-Travel:** `jj op log`, `jj undo`, `jj squash -u`, `jj split -m`, `jj absorb`.

2026 Agentic VCS Decision Rule (wavect.io & jj-vcs/jj): Hybrid git worktree + Colocated jj

- Use `git worktree` + shared cache for fast cold starts (or when Git LFS / submodules / `.gitattributes` are required); layer colocated `jj workspace` when parallel agents rewrite stacked PRs, resolve conflicts asynchronously, and need instant `jj undo` rollback.

Lab 07: Context-Firewall Subagents & Peer Council (subagent_council.py)

Walking through labs/lab_07_subagent_firewalls_and_peer_council/ step by step

Step 1 & 2: Read-Only Firewall & 3-Tier PR Triage

- **Step 1 – Subagent Context Firewall (`run_isolated_subagent`):**

1. Enforces strict tool allowlists (`code-explorer` blocks `Write/Edit/Bash`).
2. Discards the child's multi-thousand-token exploration trace and returns only a bounded ≤ 500 -token cited handoff summary.

- **Step 2 – Multi-Agent PR Triage (`synthesize_pr_triage`):**

Aggregates findings from Security, Performance, Architecture, Style, and Edge-Case reviewers, drops findings with `confidence < 80`, and buckets into `MUST_FIX`, `RECOMMENDED`, and `CONSIDER`.

Step 3 & 4: Anonymized llm-council & Worktree Check

- **Step 3 – Anonymized Borda Peer Council (`run_anonymized_council`):**

1. Scrubs model/agent names into `Candidate_A..N` before cross-review.
2. Computes Borda count scores across peer ballots while blocking self-voting manipulation, and synthesizes the winning design.

- **Step 4 – Parallel Worktree Collision Guard (`verify_worktree_isolation`):**

Blocks concurrent write subagents that share a working tree or have overlapping file globs.

- **Verify Lab 07 in $< 0.05s$:**

```
./labs/lab_07_subagent_firewalls_and_peer_council/self_diagnose.sh
```

Lab 07 Key Takeaway: Multi-Agent Architecture = Context Isolation + Bias-Free Verification

- Subagent firewalls keep the parent context window clean, while anonymized peer review and confidence-filtered PR triage catch subtle architectural defects before merge.

4-Tier “Governed by Design” (openworker), Omnigent & Hooks

Hard Floors, Earned Autonomy, Stateful Contextual Policies, and Egress-Proxy Sandboxing

Tier 1: Hard Floors

Unbypassable Gates:

- “The fixer is never the only checker.”
- `PreToolUse` hook returning **Exit Code 2** deterministically blocks edits to locked tests, secrets, or prod configs.
- No prompt or `auto` mode can talk past a Hard Floor.

Tier 2: Earned Autonomy

Stateful & Reviewer Gates:

- Independent reviewer model + `semgrep` gates auto-mode.
- **Omnigent Stateful Policies:** Escalates `git push` to human approval *after* `npm install`, and pauses after every **\$100** of LLM spend.

Tier 3: Audit Provenance

Tamper-Evident Ledger:

- Every tool invocation is logged with `user-approved`, `auto-approved`, or `denied`.
- Records exact command hash, reviewer rationale, and git tree state.
- Enables full post-incident replay.

Tier 4: Egress Sandbox

OpenShell & Omnigent:

- Agent commands run in a local/cloud container (`Modal/Daytona`).
- **Egress-Proxy Token Injection:** The agent *never* sees raw GitHub tokens; tokens are injected only in the egress proxy!

Deterministic Lifecycle Hooks Across `claude` (`.claude/settings.json`), `agy` (`.agents/plugins/`) & Omnigent

- **`PreToolUse` (Feedforward Gate):** Inspects tool JSON on stdin; `exit 0` allows, **`exit 2` blocks** and feeds `stderr` back to the agent.
- **`PostToolUse` (Feedback Sensor):** Runs auto-formatters and remediation linters immediately after every `Write/Edit`. • **`Stop`:** Blocks turn completion until tests pass.

Remediation-Aware Architectural Linters (OpenAI Harness Engineering)

Why standard linter messages fail agents, and how to emit [VIOLATION] → [WHY] → [HOW TO FIX] recipes

Why Traditional Human Linter Messages Fail Agents

- **Cryptic Error Codes:** A standard linter outputs:

```
src/api/orders.py:42: E901 Forbidden import db.session
```

- **What the Coding Agent Does Next:**

1. Does not know *which* module it is supposed to call instead.
2. Wastes 4 tool calls grepping the codebase for workarounds, or worse, adds `# noqa / pylint: disable` to silence the check!
3. Pollutes the context window with trial-and-error guesses.

OpenAI's Remediation-Aware Linter Pattern

- **Write Linter Errors for an LLM Reader ([VIOLATION] → [WHY] → [HOW TO FIX]):**

```
[VIOLATION] Direct import of 'db.session' in 'src/api/orders.py:42' violates Layer Boundary Rule ARCH-03 (API ↗ DB).
```

```
[WHY] Route handlers must remain stateless and testable without a live database connection.
```

```
[HOW TO FIX] Replace 'from db.session import query' with 'from services.order_service import OrderRepository' and inject via 'Depends(get_order_repo)'.
```

- Enables **single-turn autonomous self-repair** on the very next step!

Encoding Architectural Invariants as Fast AST / Regex Structural Checks

- Enforce layer dependency DAGs (`ui → api → service → repo`), ban raw floats in monetary modules, require structured logging, and forbid `# noqa` suppressions via deterministic `PostToolUse` checks.

Bounded Autonomous Loops: Karpathy's autoresearch & program.md

The 3-File Harness Separation, The Simplicity Criterion, and the Automated Git Keep-or-Revert Ratchet

1. prepare.py (Immutable Sensor)

- **Read-Only Ground Truth:** Contains dataset prep, fixed 5-minute wall-clock budget, and `evaluate_bpb()`.
- Protected by `PreToolUse` Exit-Code-2 Hard Floor: the agent is **strictly forbidden from editing the evaluator** or installing packages.

2. train.py (Bounded Action Space)

- **Single Mutable File:** The *only* file the autonomous agent is permitted to edit (model architecture, optimizer, schedule).
- Eliminates cross-repo blast radius so every experiment is a clean, self-contained diff.

3. program.md ("Research Org Code")

- **Human-Authored Harness Spec:** The Markdown instructions programming the autonomous research loop.
- Enforces output backpressure (`> run.log 2>&1`) and the **Simplicity Criterion**.


The Automated Git Keep-or-Revert Ratchet (~12 Experiments/Hr, ~100 Overnight) & The Simplicity Criterion

- **Closed-Loop Ratchet:** 1. `Propose Hypothesis` → 2. `Edit train.py & git commit` → 3. `Run 5-min budget (> run.log 2>&1)` → 4. `Parse val_bpb & log to untracked results.tsv` → **If improved & passes Simplicity Criterion: KEEP commit (ratchet advances)** → **If worse or crash: git reset -hard HEAD^ (instant revert)!**
- **Karpathy's Simplicity Criterion:** *"All else being equal, simpler is better. A 0.001 improvement that adds 20 lines of hacky code? Revert. A 0.001 improvement from deleting code? Definitely keep!"*

Why “Lights-Off” Software Factories Fail (Ib5GBkD555M): The RLVR Horizon Gap

Why optimizing solely for $t = 0$ test pass causes Week-8 architectural collapse, and the 4-Phase Front-Loaded Fix

Feature Velocity & Maintainability Over Time (Weeks 1–12)



The RLVR Reward Horizon Gap (Ib5GBkD555M)

- **Why Dark Factories Collapse at Week 8:** RLVR trains models to maximize $t = 0$ **immediate rewards** (passing current tests, compiling cleanly).
- But software architecture lives on a $t = 6$ **month horizon**! Left unattended in a “lights-off” factory, agents duplicate abstractions, copy-paste 400-line handlers, and create spaghetti coupling that passes $t = 0$ tests until the codebase collapses.
- **Deterministic Harnesses Raise the Floor, Not the Ceiling:**
Linters and tests prevent syntax/layer bugs (floor), but **humans must own the architectural ceiling** via 4-Phase Front-Loaded Alignment.

Keeping the “Lights On”: Dex Horthy’s 4-Phase Front-Loaded Alignment

- **Phase 1: Product** (User value) → **Phase 2: Architecture** (System boundaries) → **Phase 3: Program Design** (Exact types, interfaces, file-trees, and call-stacks co-designed with the human!) → **Phase 4: Vertical Slices** (Bounded agent implementation).

Durable Execution (Temporal), Meta-Harnesses (Omnigent) & 52x RSI

Surviving 12–16 hr runs with Temporal/DBOS, orchestrating via Omnigent, and Out-of-Tree Holdouts

1. Durable Agent Execution

Temporal, LangGraph, DBOS, Restate, Inngest:

- Why 12–16 hr loops crash in raw Python: 429 rate limits, MCP timeouts, or worker restarts at turn 18 wipe state and burn tokens.
- **Temporal Pattern:** Deterministic Agent Loop (Workflow) + non-deterministic LLM/MCP calls (Activities) → **Event-History Replay** & zero-compute **HITL Signals!**

2. Meta-Harness (Omnigent)

Databricks omnigent.ai (Zaharia et al.):

- **Combine:** Compose heterogeneous harnesses (Claude Code, Codex, agy, Pi) in YAML with 1-line swaps (Harvey advisor, Genie routing).
- **Control & Share:** Stateful policies, egress-proxy token injection, and live multiplayer URL sharing + GEPA/MemEx/RLM.

3. 52x RSI & Holdouts

Anthropic RSI, Holdouts & Amdahl’s Law:

- **52x Kernel Speedup:** Bounded single-kernel action space + cycle timer achieved **52x speedup in 8 hrs** (vs. 4x human baseline).
- **Out-of-Tree Holdouts & Amdahl’s Law:** Grade commits against isolated holdout suites; human leverage shifts to Program Design & Verification.

The Complete 2026 Production Stack: Program Design + Durable Loops + Meta-Harness Governance

- **Durable Engines (Temporal/DBOS):** Guarantee fault-tolerant state replay across multi-hour runs.
- **Meta-Harnesses (Omnigent):** Combine specialized harnesses, enforce stateful egress-proxy security, and share live multiplayer sessions.

Lab 08 (autoresearch_harness.py) & Full Course Verification (./self_diagnose_all.sh)

Walking through Lab 08's 4-Tier Governance & Keep-or-Revert Ratchet and verifying all 8 labs in < 1s

Lab 08 Step-by-Step (autoresearch_harness.py)

- **1. PreToolUse Exit-Code-2 Hard Floor (evaluate_pre_tool_use):**

Blocks edits to immutable evaluators (prepare.py, holdout_scenarios/) and logs tamper-evident provenance records.

- **2. Remediation Linter (run_remediation_linter):**

Detects layer boundary and float-money violations and formats [VIOLATION] → [WHY] → [HOW TO FIX] instructions.

- **3. Program Design & Simplicity Criterion (evaluate_ratchet_step):**

Weighs metric improvement against net LOC added and Out-of-Tree Holdout pass rate to decide KEEP_COMMIT vs. GIT_REVERT.

Master Course Verification (./self_diagnose_all.sh)

- **Run All 8 Labs + Dual-Track Suites in < 1s:**

./self_diagnose_all.sh

- **Phase 1 Labs (01–04):** prompt_steerer.py · wire_trace_inspector.py · reppit_orchestrator.py · context_compactor.py — **PASS**

- **Phase 2 Labs (05–08):** context_compiler.py · skill_and_mcp_harness.py · subagent_council.py · autoresearch_harness.py — **PASS**

- **Next Step in the Curriculum Track:** For formal multi-service brownfield specifications, continue to the companion Spec-Driven Development course.

Course Complete! Thank You — Alejandro Kaspar, AI Forward Deployed Engineer (CC BY 4.0 / Apache 2.0)

- Clone the repository, switch between the agy (Google Antigravity) and claude (Anthropic Claude Code) branches, and apply Phase 1 Vibe Coding and Phase 2 Harness Engineering to your production systems!